

Income-Related Performance Gaps in Vision-Language Models

Diya Tiwari

Riya Katikar

Abstract

We study whether pretrained vision-language models perform differently on pictures of household objects from four income groups in Dollar Street. The benchmark contains 168 images from six categories, with seven images in every category-income cell. We compare CLIP, BLIP, Qwen2.5-VL, InternVL3.5, and YOLO-World without fine-tuning. CLIP classification and every automatic caption measure scored higher on the highest-income quartile than the lowest, while Qwen and InternVL classification changed little. A two-rater review showed that vocabulary matching understated BLIP caption quality. The measured gap depends on the model, task, and evaluation metric.

1. Introduction

The same household object can look different across homes. A stove may be built into a kitchen, placed on a table, or use solid fuel. Models trained largely on web image-text pairs may be more familiar with some forms than others. Earlier Dollar Street work found income-related differences in CLIP alignment and retrieval [4]. We ask: How does performance on six household-object categories vary across four income quartiles when each category is equally represented? Income groups the images; it is not treated as the cause of an error.

2. Method

Dataset. We used the public 1,600-row Dollar Street test split [6]. Income records monthly consumption per adult equivalent in PPP-adjusted US dollars. After removing 45 multi-label rows, we formed quartiles from 1,555 eligible records. The evaluation subset has 168 images from 44 countries and six categories: roof, light source, stove, trash container, switch, and footwear. Each of the 24 category-quartile cells has seven images. This controls category frequency, although the subset is not representative of the population and regional counts remain unequal.

Models and tasks. CLIP ViT-B/32 [5] performed six-way zero-shot classification with “a photo of a {label}.” BLIP-base [3] generated an unprompted caption and a caption beginning with the label-free prefix “the main household object in this image is.” Qwen2.5-VL-3B-Instruct [1] and InternVL3.5-2B [7] used the same forced-choice classification instruction and short, location-neutral caption instruction. YOLO-World-small-v2 [2] received the six labels with a 0.25 confidence threshold. Dollar Street has no bounding boxes, so YOLO-World is reported as an image-level correct-class detection rate; mAP and IoU cannot be calculated.

CLIP used its 224-pixel checkpoint processor and BLIP its 384-pixel processor. Qwen used a 256-token visual budget, InternVL used 448×448 RGB inputs with ImageNet normalization, and YOLO-World used the Ultralytics pipeline. Generation was deterministic. The repository includes the manifest, identifiers, prompts, outputs, and run settings.

Evaluation. Classification uses top-1 accuracy. Caption recall checks for a predefined category term or synonym using whole-term matching. Two raters independently reviewed all 336 BLIP captions for object correctness, non-informativeness, and hallucination, then resolved disagreements. We report raw agreement and Cohen’s κ . A fixed 84-image sample checked Qwen captions. The main gap is Q4 minus Q1, with a 95% category-stratified bootstrap interval from 2,000 resamples using seed 2026. We checked all 1,344 prediction rows for coverage, duplicate keys, and missing metrics.

3. Results and analysis

The classifiers had similar overall accuracy (86.3–86.9%), but their quartile results differed (Table 1). CLIP’s Q4–Q1 gap was 16.7 points (95% CI 2.4 to 28.6). The intervals for Qwen’s 2.4-point gap (−7.1 to 14.3) and InternVL’s zero-point gap (−9.5 to 9.5) included zero. Thus, the CLIP pattern did not appear in both instruction-following classifiers.

Automatic caption recall was higher in Q4 than Q1 for every caption condition. The gaps were 47.6 points

Table 1. Scores (%) by income quartile. Each quartile has 42 images. Metrics are task-specific and should not be compared as if they measured the same output.

Model and task	Q1	Q2	Q3	Q4	All
CLIP classification	76.2	81.0	95.2	92.9	86.3
Qwen classification	85.7	88.1	83.3	88.1	86.3
InternVL classification	88.1	92.9	78.6	88.1	86.9
YOLO-World detection	11.9	14.3	14.3	21.4	15.5
BLIP caption, baseline	28.6	52.4	54.8	76.2	53.0
BLIP caption, prompted	26.2	54.8	45.2	78.6	51.2
Qwen caption	35.7	64.3	71.4	81.0	63.1
InternVL caption	42.9	59.5	66.7	64.3	58.3

for unprompted BLIP, 52.4 for prompted BLIP, 45.2 for Qwen, and 21.4 for InternVL. Manual review gave a different view of BLIP. Object correctness was 80.4% without the prompt and 83.9% with it, compared with automatic recall of 53.0% and 51.2%. The prompt reduced automatic recall by 1.8 points, but improved human-rated correctness by 3.6 points and reduced hallucinations from 16.7% to 12.5%. Agreement was 89.3% ($\kappa = 0.65$) for object correctness and 94.0% ($\kappa = 0.77$) for hallucination. The term list missed reasonable descriptions such as bucket or basket for a trash container.

Category results showed different weaknesses. CLIP was weakest on light sources (60.7%); Qwen and InternVL were weakest on roofs (39.3% and 50.0%). YOLO-World found 57.1% of footwear and 25.0% of stoves, but none of the roofs, light sources, or switches. In one failure case, CLIP labelled a small red fuel can as a stove while BLIP described the visible can. In another, a roof occupied little of the frame and was classified as a light source. These cases explain why we include category results and manual caption review.

4. Discussion and limitations

The results do not support one common income pattern across all models. CLIP and the automatic caption measures favored Q4, whereas Qwen and InternVL classification were nearly flat. The BLIP comparison also depended on the metric because term recall missed wording that raters accepted.

The study covers 168 selected images and six categories. Balancing removes category-frequency differences but does not control country, region, image quality, scene complexity, or variation within a category. Quartiles are defined within the test split, and the intervals describe this benchmark only. Models never received income information, so the results show an association and do not establish that income caused errors.

The Gradio interface demonstrates the five models on uploaded images; it is not a second evaluation dataset.

5. Related work

Dollar Street provides household-object images with country, region, and income metadata [6]. Nwatu et al. evaluated CLIP alignment and prompt-based retrieval across income groups [4]. We use a category-balanced subset, test six-way recognition and captioning, and compare five model families on the same images. CLIP links the experiment to earlier work [5], while BLIP adds caption generation [3]. Qwen2.5-VL and InternVL3.5 add instruction-following VLMs [1, 7], and YOLO-World adds open-vocabulary detection [2]. A gap observed with one model or automatic metric did not always carry over to another.

References

- [1] Shuai Bai et al. Qwen2.5-vl technical report. arXiv preprint arXiv:2502.13923, 2025. 1, 2
- [2] Tianheng Cheng et al. Yolo-world: Real-time open-vocabulary object detection. In CVPR, pages 16901–16911, 2024. 1, 2
- [3] Junnan Li, Dongxu Li, Caiming Xiong, and Steven Hoi. Blip: Bootstrapping language-image pre-training for unified vision-language understanding and generation. In ICML, pages 12888–12900, 2022. 1, 2
- [4] John Nwatu, Oana Ignat, and Rada Mihalcea. Bridging the digital divide: Performance variation across socioeconomic factors in vision-language models. In EMNLP, pages 10686–10702, 2023. 1, 2
- [5] Alec Radford et al. Learning transferable visual models from natural language supervision. In ICML, pages 8748–8763, 2021. 1, 2
- [6] William A. Gaviria Rojas et al. The dollar street dataset: Images representing the geographic and socioeconomic diversity of the world. In NeurIPS Datasets and Benchmarks, 2022. 1, 2
- [7] Weiyun Wang et al. Internvl3.5: Advancing open-source multimodal models. arXiv preprint arXiv:2508.18265, 2025. 1, 2